

Understanding the Art of Voting: A Visual Approach

Experimentation & Verification on the Performance of a Voting Classifier

Motivation

Ensemble method can also be known as the wisdom of the crowd, in which a conclusion is made based on answers given by multiple parties. A classifier for an ensemble method can be made from multiple different machine learning algorithms of the user's liking. One such classifier, known as soft voting classifier, first averages out the prediction probability from each individual classifier before predicting the class based on the highest probability. It is believed using a voting classifier could outperform the classifiers constituting it. This is due to the law of large numbers, in which having many bias classifiers predicting a class will converge the prediction to that class, in other words, the probability of predicting a majority class is higher. The downside is if the majority predicted class is wrong, error is propagated to the voting classifier. If each classifier has diverse performance and error, the same error will not compound in the voting classifier and shouldn't deteriorate the voting classifier's performance. Hence, in this project, experiment is done to evaluate the performance of a soft voting classifier.

Methodology

The machine learning task to work on would be a binary classification task. Three completely different machine learning algorithms for diversity, namely logistic regression (LR), decision trees (DT) and multi-layer perceptron (MLP) will be trained separately and compared with the soft voting classifier (VC), which is made up of those three machine learning algorithms. The dataset to be used will be the [chronic kidney disease](#) dataset from the UCI Machine Learning Repository. The experiment is outlined as follows. First, pre-process the dataset by only using numerical attributes. The dataset is then split randomly for 60% training data and 40% testing data. For the visualization of the classification boundary, only two attributes will be used and they are the top 2 most important attributes shared commonly among each algorithm using Recursive Feature Elimination. Each classifier is fitted with the training dataset and the training accuracy is obtained which is compared with its test accuracy in a bar chart. A contour plot describing the classification regions for each classifier and where those test data lie at using the two attributes from the dataset is plotted for analysis.

Results

Figure 1 shows the bar chart comparing the training and testing accuracy for each classifier.

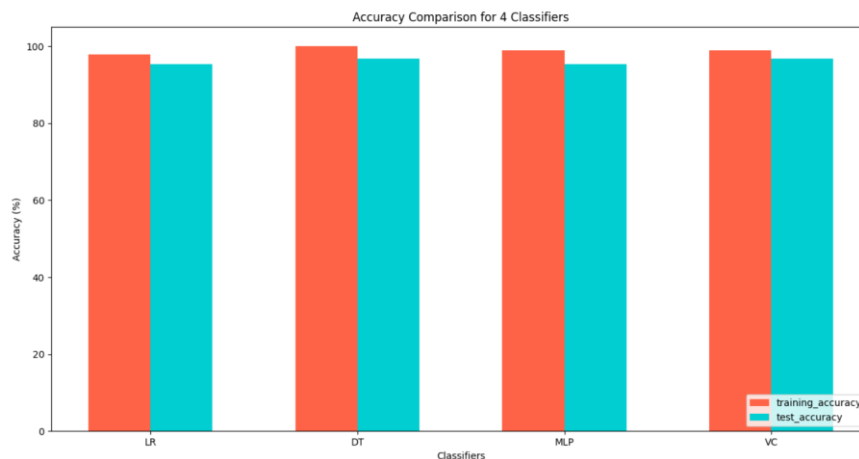


Figure 1: Accuracy Comparison for the Four Classifiers.

Figure 2 shows the classification region or the decision boundary of each classifier. The data points here are from the test set (40% of the whole dataset).

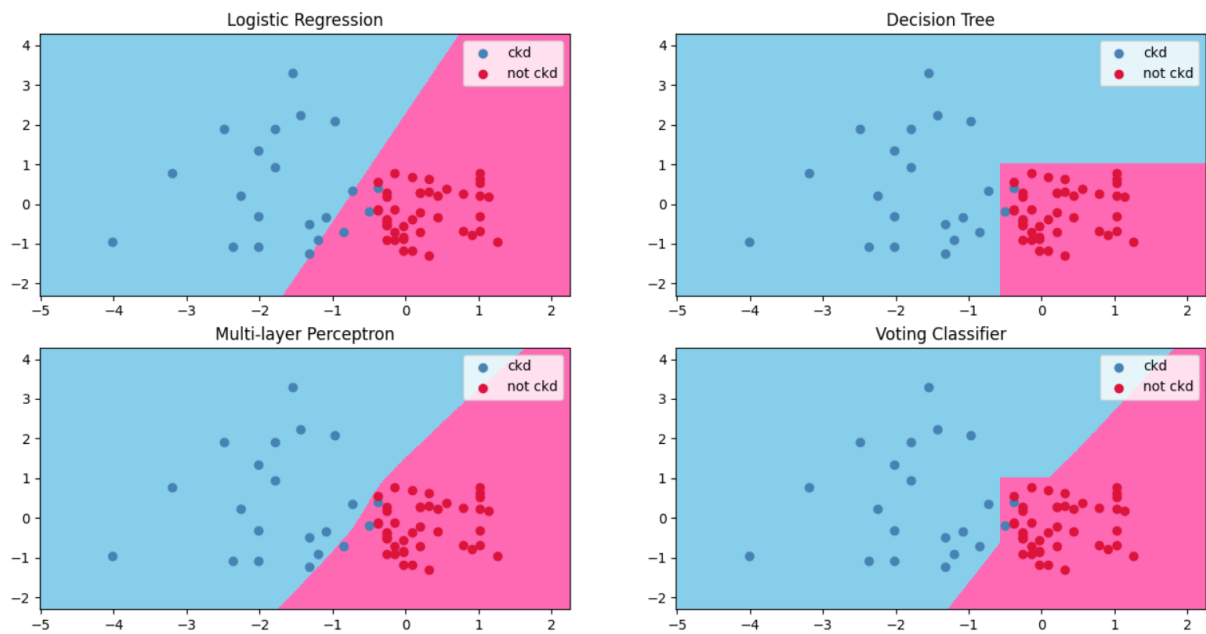


Figure 2: Classification region for all four classifiers. Chronic kidney disease (ckd) is the blue region whereas the opposite class (not ckd) is the pink region.

Table 1 explicitly states the training and testing accuracy values and their differences for each classifier.

Table 1: Comparison of training and test accuracies with their differences for each classifier

Classifier	Training Accuracy (%)	Test Accuracy (%)	Difference (%)
Logistic Regression	97.872	95.312	2.560
Decision Tree	100.00	96.875	3.125
Multi-layer Perceptron	98.936	95.313	3.623
Soft Voting Classifier	98.936	96.875	2.061

Based on Figure 2, it can be seen that the three non-fusion classifiers all have their own decision boundary, in which LR has a linear classification region whereas DT and MLP have a non-linear classification region. Due to the linearity of the LR, some of the ckd samples are predicted incorrectly. Despite that, MLP, which is non-linear, still incorrectly predicted some of the test samples, making it performed similarly to LR. On the other hand, the box-shaped boundary region of DT allowed DT to prevent misclassifying the test samples like LR and MLP. As for VC, it is noted that VC contains the characteristics of DT and MLP. However, the stretched-out classification region of the VC differed in a way that this region is slightly shifted to the right, making the VC not misclassifying the test sample like that in MLP. Based on each classifier's classification region, no traces of over-/under-fitting can be observed. It can be inferred that the performance of VC is highly dependent on the classifiers constituting it. If those classifiers are similar in performance and classification region, the VC might not perform exceptionally well. For example, during training, DT scored perfectly but not for LR and MLP, and this error was propagated to the VC, as seen in Table 1 whereby the VC's training accuracy isn't perfect, but better than LR. Therefore, VC might be able to perform remarkably well than the individual classifiers constituting it if there are sufficient number of those classifiers with good performances and diverse decision boundaries.

Appendix : Source Code

```
1 import matplotlib.pyplot as plt
2 from matplotlib.colors import ListedColormap
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.feature_selection import RFE
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.tree import DecisionTreeClassifier
8 from sklearn.neural_network import MLPClassifier
9 from sklearn.ensemble import VotingClassifier
10 import numpy as np
11 import pandas as pd
12
13
14 # Loads the dataset given the path and the filename
15 def load(filepath):
16     data = []
17     f = open(filepath, "r")
18     lines = f.read().split("\n") # split the dataset according to the lines/rows
19     for line in lines:
20         data.append(line.split(",")) # split a line at each comma
21     f.close()
22     return data
23
24
25 # Filter out non-numerical attributes from the dataset and process the class labels
26 def filter_numerical_attributes_and_class_label(data):
27     numerical_col_ind = [0, 1, 9, 10, 11, 12, 13, 14, 15, 16, 17] # all the 11 numeric attributes column index
28     df = pd.DataFrame(data)
29     class_label = df.iloc[:, 24]
30     class_label = class_label.replace("ckd", 0)
31     class_label = class_label.replace("notckd", 1)
32     df = df.iloc[:, numerical_col_ind]
33     return df.to_numpy(dtype=float), class_label.to_numpy()
34
35
36 # For analysis purpose, get the index of rows with missing data. Missing data is represented as '?' in the columns
37 def get_index_of_missing_data(data):
38     ind_list = []
39     for idx, d in enumerate(data):
40         if '?' in d:
41             ind_list.append(idx)
42     return ind_list
43
44
45 # remove rows with missing data
46 def remove_rows_of_missing_data(data):
47     filtered = []
48     for d in data:
49         if '?' not in d:
50             filtered.append(d)
51     return filtered
52
53
```

```

94 # simple function to calculate the accuracy using classifiers and datasets
95 def get_classifier_accuracy(classifier, x, y):
96     # get predictions
97     y_predict = classifier.predict(x)
98     accuracy = 0.0
99     for index, actual in enumerate(y):
100         if y_predict[index] == actual:
101             accuracy += 1.0
102     return accuracy / float(len(y)) * 100.0
103
104
105 # function to plot the classification region for any classifier
106 def plot_clf_region(classifier, x, y):
107     # first get the bounds of the domain for the contour plot, this would be the max and min from the two attributes
108     min1, max1 = x[:, 0].min() - 1, x[:, 0].max() + 1
109     min2, max2 = x[:, 1].min() - 1, x[:, 1].max() + 1
110
111     # create ranges of data for attributes x1 and x2
112     res = 0.01 # resolution of the contour plot
113     x1 = np.arange(min1, max1, res)
114     x2 = np.arange(min2, max2, res)
115
116     # forming a grid, basically means one row of x1 for each point on the y-axis and one column of x2 for each point
117     # on the x-axis
118     xx, yy = np.meshgrid(x1, x2)
119
120     # flatten the grid to allow the samples from it to be used for prediction by the classifier
121     r1, r2 = xx.flatten(), yy.flatten()
122     r1, r2 = r1.reshape((len(r1), 1)), r2.reshape((len(r2), 1))
123
124     # stack the r1 and r2 vectors horizontally to create [attribute1, attribute2] format for classifier use
125     grid = np.hstack((r1, r2))
126
127     # use the classifier to predict all the points in the grid
128     y_predict = classifier.predict(grid)
129
130     # reshape the prediction vector back to mesh-grid format for plotting purposes
131     z = y_predict.reshape(xx.shape)
132
133     # finally plot it using matplotlib's filled contour plot
134     cmap_contour = ListedColormap(['skyblue', 'hotpink'])
135     plt.contourf(xx, yy, z, cmap=cmap_contour)
136
137     # also add the test data points to it
138     label = ['ckd', 'not ckd']
139     c_scatter = ['steelblue', 'crimson']
140     for class_label in range(2):
141         # get the rows of the test data that is of class_label 0 or 1
142         row = np.where(y == class_label)
143         plt.scatter(x[row, 0], x[row, 1], label=label[class_label], c=c_scatter[class_label])
144     plt.legend(loc="best")
145
146
147 # dataset used is the chronic kidney disease classification dataset from UCI Machine Learning Repository
148 # this dataset has 11 numerical attributes and 13 nominal attributes (24 attributes in total, excluding class label)
149 # path that leads to the location of the dataset (stored locally in own device)
150 pathname = "Chronic_Kidney_Disease/chronic_kidney_disease_data.txt"
151

```

```

112 # load the dataset into a list
113 dataset = load(pathname)
114 print("Visualize the first 5 data")
115 print(dataset[0:5])
116 print("\nNumber of data : ", len(dataset)) # there would be 400 number of data in total
117 # get the indices of rows with missing data, a row has missing data if it has a '?' in at least one column
118 row_miss = get_index_of_missing_data(dataset)
119 print("\nNumber of data with missing columns : ", len(row_miss)) # there would be 242 number of unusable data
120 dataset = remove_rows_of_missing_data(dataset)
121 print("\nNumber of data after filtering out missing columns : ", len(dataset)) # only 158 number of useful data
122
123 # since in this project, only numerical attributes are used. Filter out non-numerical attributes
124 # also separate and process the class labels, ckd and notckd as 0 and 1 respectively
125 X, Y = filter_numerical_attributes_and_class_label(dataset)
126 x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.4, random_state=40)
127
128 # Scale the data attributes to their unit variance
129 scale = StandardScaler()
130 x_train = scale.fit_transform(x_train)
131 x_test = scale.transform(x_test)
132
133 # Initialize each classifier: Logistic regression, decision tree, multi-layer perceptron and voting classifier
134 lr = LogisticRegression(random_state=10)
135 dt = DecisionTreeClassifier(random_state=12)
136 mlp = MLPClassifier(random_state=30)
137 vc = VotingClassifier(estimators=[('lr', lr), ('dt', dt), ('mlp', mlp)], voting='soft')
138
139 # Perform feature selection on the dataset for each classifier using Recursive Feature Selector
140 # Unfortunately, MLPClassifier isn't supported by RFE function. Instead of changing the classifier used,
141 # only the logistic regression and the decision tree classifier is used to obtain the 2 important attributes
142 n_features = 4
143 selector_lr = RFE(estimator=lr, n_features_to_select=n_features, step=1) # for logistic regression
144 selector_lr = selector_lr.fit(x_train, y_train)
145 selector_dt = RFE(estimator=dt, n_features_to_select=n_features, step=1) # for decision tree classifier
146 selector_dt = selector_dt.fit(x_train, y_train)
147
148 # Visualize and look for the 2 common important features for all classifiers
149 print("Ranking of Features for LR : ", selector_lr.ranking_)
150 print("Ranking of Features for DT : ", selector_dt.ranking_)
151
152 # as the classifiers differ immensely, selecting the two attributes is based on own judgement
153 # from own perspective, it seems col 8 and col 9 will be the 2 attributes chosen to proceed with the experiment
154 # changing the random state of each classifier and the train_test_split func might alter this decision!
155 feature_to_keep_idx = [8, 9]
156
157 # Reduce the dataset to only two of the selected attributes
158 new_x_train = x_train[:, feature_to_keep_idx] # x_train now only contains the two selected feature
159 new_x_test = x_test[:, feature_to_keep_idx] # x_test now only contains the two selected feature
160
161 # Train each classifier
162 lr.fit(new_x_train, y_train)
163 dt.fit(new_x_train, y_train)
164 mlp.fit(new_x_train, y_train)
165 vc.fit(new_x_train, y_train)
166
167 # Retrieve their training accuracy
168 train_acc_lr = get_classifier_accuracy(lr, new_x_train, y_train)
169 train_acc_dt = get_classifier_accuracy(dt, new_x_train, y_train)
170 train_acc_mlp = get_classifier_accuracy(mlp, new_x_train, y_train)
171 train_acc_vc = get_classifier_accuracy(vc, new_x_train, y_train)
172 print("Training Accuracy for Logistic Regression : ", train_acc_lr)

```

```

173 print("Training Accuracy for Decision Tree : ", train_acc_dt)
174 print("Training Accuracy for Multilayer Perceptron : ", train_acc_mlp)
175 print("Training Accuracy for Voting Classifier : ", train_acc_vc)
176
177 # Use the classifiers to make predictions on test set and get the test accuracy
178 test_acc_lr = get_classifier_accuracy(lr, new_x_test, y_test)
179 test_acc_dt = get_classifier_accuracy(dt, new_x_test, y_test)
180 test_acc_mlp = get_classifier_accuracy(mlp, new_x_test, y_test)
181 test_acc_vc = get_classifier_accuracy(vc, new_x_test, y_test)
182 print("\nTest Accuracy for Logistic Regression : ", test_acc_lr)
183 print("Test Accuracy for Decision Tree : ", test_acc_dt)
184 print("Test Accuracy for Multilayer Perceptron : ", test_acc_mlp)
185 print("Test Accuracy for Voting Classifier : ", test_acc_vc)
186
187 # Comparison of each classifier with their training and test accuracy using a bar chart
188 train_acc = np.array([train_acc_lr, train_acc_dt, train_acc_mlp, train_acc_vc]) # height of training accuracy
189 test_acc = np.array([test_acc_lr, test_acc_dt, test_acc_mlp, test_acc_vc]) # height of test accuracy
190 bar_width = 0.3 # the width of each bar in the bar chart
191 number_of_bars = 4 # since we have 4 classifiers
192 ind = np.arange(number_of_bars) # index position for the bar chart
193
194 plt.bar(ind, train_acc, color='tomato', width=bar_width, label='training_accuracy')
195 plt.bar(ind + bar_width, test_acc, color='darkturquoise', width=bar_width, label='test_accuracy')
196 plt.xlabel('Classifiers')
197 plt.ylabel('Accuracy (%)')
198 plt.title('Accuracy Comparison for 4 Classifiers')
199 plt.xticks(ind + bar_width / 2, ('LR', 'DT', 'MLP', 'VC'))
200 plt.legend(loc='lower right')
201 plt.show()
202
203 # Visualization of each classifier's classification boundary using a contour plot
204 plt.subplot(2, 2, 1)
205 plot_clf_region(lr, new_x_test, y_test)
206 plt.title("Logistic Regression")
207 plt.subplot(2, 2, 2)
208 plot_clf_region(dt, new_x_test, y_test)
209 plt.title("Decision Tree")
210 plt.subplot(2, 2, 3)
211 plot_clf_region(mlp, new_x_test, y_test)
212 plt.title("Multi-layer Perceptron")
213 plt.subplot(2, 2, 4)
214 plot_clf_region(vc, new_x_test, y_test)
215 plt.title("Voting Classifier")
216 plt.show()

```